

SSE3052: Embedded Systems Practice

Jinkyu Jeong

jinkyu@skku.edu

Computer Systems Laboratory

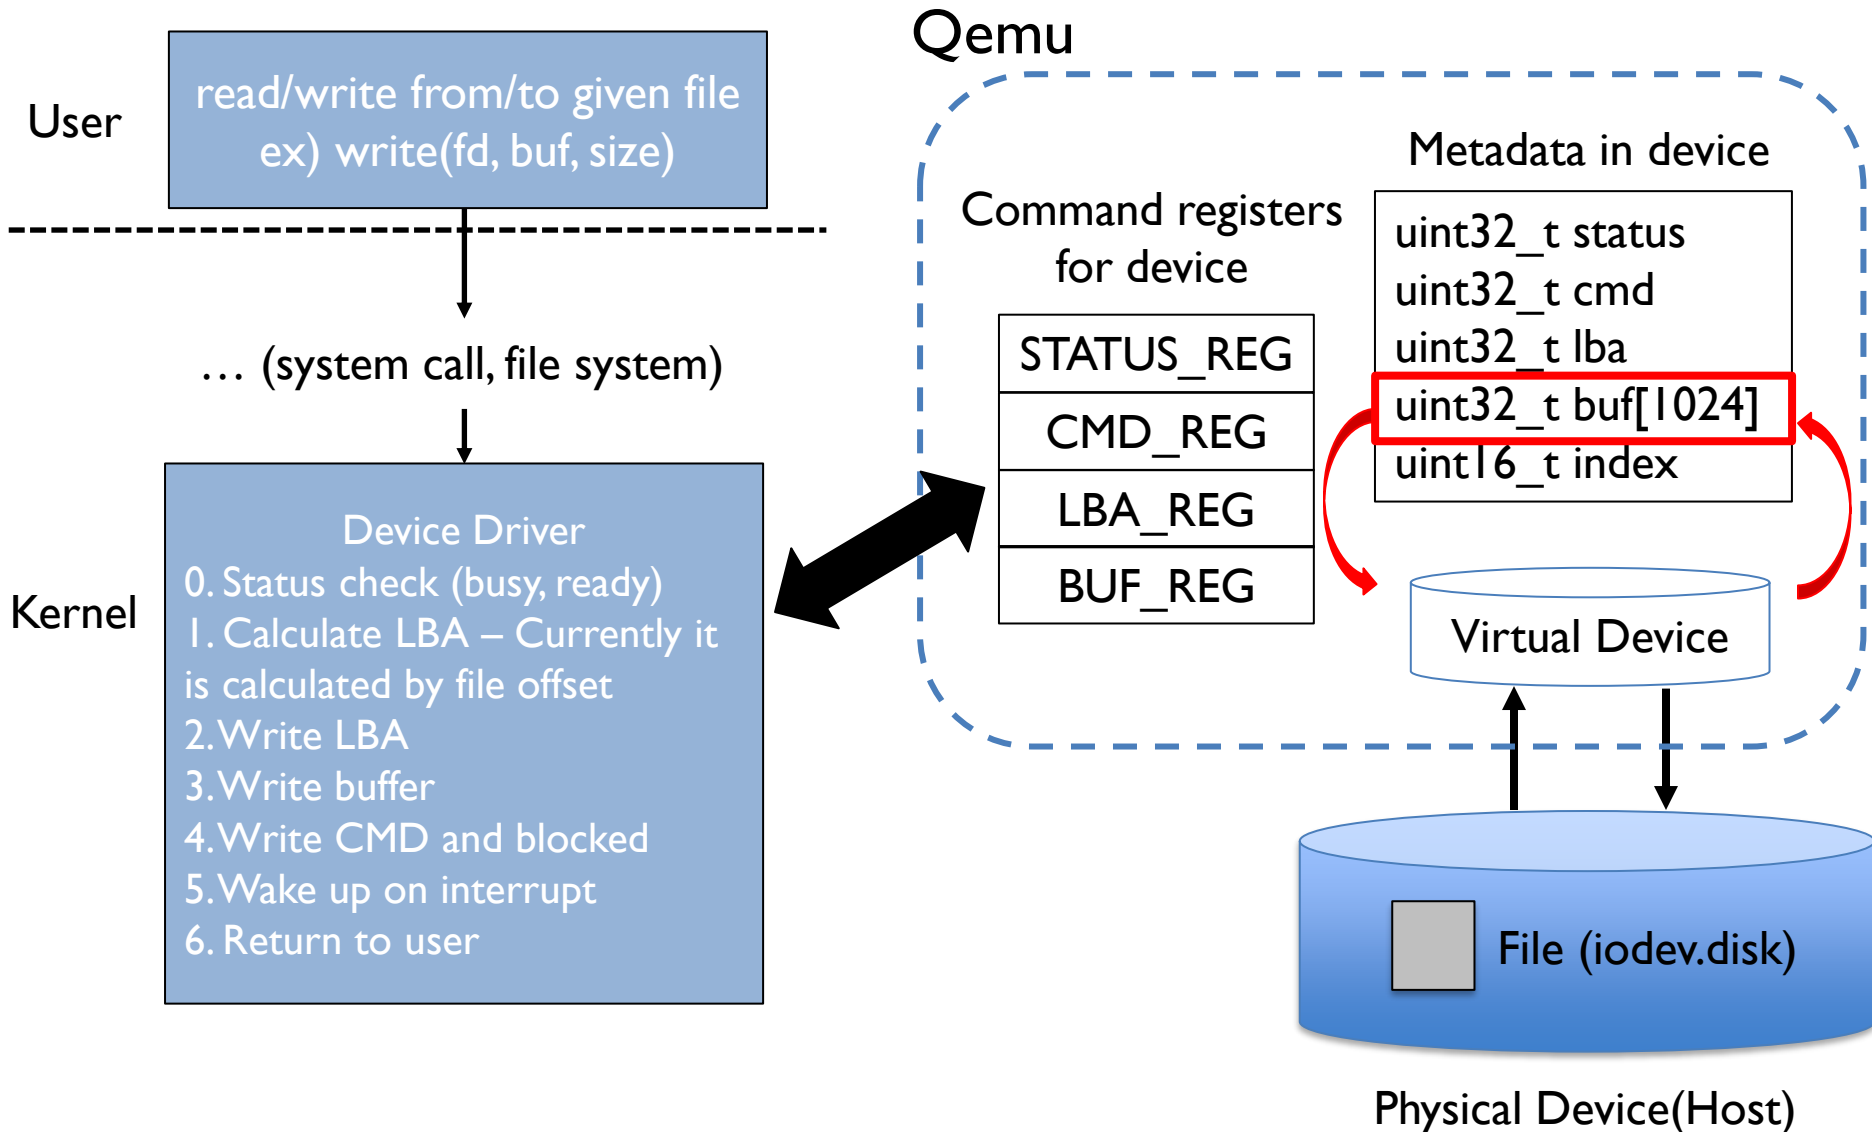
Sungkyunkwan University

<http://csl.skku.edu>

Today's goal

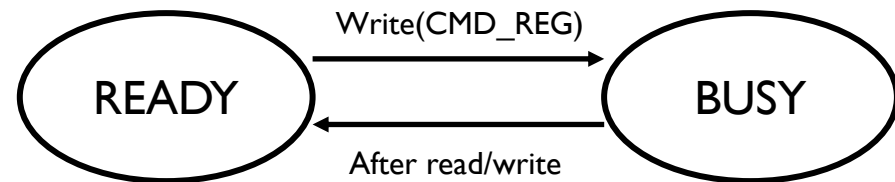
- Add new soft block device (no skin)
- Implement interrupt based IO

Overview



The New Soft Device

- 4KB read/write block device
- Four device registers
 - STATUS_REG (read-only)
 - DEV_BUSY: device busy
 - DEV_READY: ready to receive a new command
 - CMD_REG (write-only)
 - CMD_READ / CMD_WRITE
 - LBA_REG (write-only)
 - Logical block number in 4KB
 - BUF_REG (read/write): 4KB buffer
 - 1024 * 4B writes fill in the buffer
 - 1024 * 4B reads fetch data from the buffer



Protocols

- Description of interface btw. goldfish and device
- Following two slides are protocols of read/write
- Template code of QEMU and device driver are opened on iCampus

Write Protocol

1. Buffer copy from user
2. Status check
 - Busy-wait for device's state become DEV_READY
3. Write LBA
4. Fill in the device buffer
 - $1024 * 4B$ writes to BUF_REG
5. Write command (CMD_WRITE)
6. Wait for interrupt

Read Protocol

1. Status check
 - Busy-wait for device's state become DEV_READY
2. Write LBA
3. Write command (CMD_READ)
4. Wait for interrupt
5. Read device buffer
 - 1024 * 4B reads to BUF_REG
6. Buffer copy to user

QEMU Protocol

- STATUS_REG (read-only)
 - Return current state of device (CMD_READY, CMD_BUSY)
- CMD_REG (write-only)
 - Switch current state to CMD_BUSY
 - Control written operation (READ_CMD, WRITE_CMD)
 - Value of lba is used
 - Raise interrupt to host
- LBA_REG (write-only)
 - Assign written value to lba
- BUF_REG (read/write)
 - Write: assign written value to device buffer
 - Read: return value of device buffer

Modify your QEMU

- Refer to “qemu_devices_with_led.patch” (From resources on week 4)
 - Make new file for emulate new soft device
 - e.g., hw/android/goldfish/iodev.c
 - Edit Makefile.qemu l-target.mk
 - To compile hw/android/goldfish/iodev.c
 - Call initialization of new device in hw/i386/pc.c
 - goldfish_iodev_init(void)
 - Include initialization in include/hw/android/goldfish/device.h
 - void goldfish_iodev_init(void)

Modify your Goldfish

- Refer to “goldfish_kernel_with_led.patch” (From resources on week 4)
 - Add kernel configuration for new soft device
 - `drivers/misc/Kconfig`
 - Edit Makefile
 - `drivers/misc/Makefile`
 - Write device driver for new soft device
 - e.g., `drivers/misc/goldfish_iodev.c`

(Goldfish) goldfish_iodev_probe (I)

```
static int goldfish_iodev_probe(struct platform_device *pdev)
{
    int ret;
    struct resource *r;
    int error;

    printk(KERN_ERR "iodev probe started\n");
    data = devm_kzalloc(&pdev->dev, sizeof(*data), GFP_KERNEL);
    if (data == NULL)
        return -ENOMEM;

    r = platform_get_resource(pdev, IORESOURCE_MEM, 0);
    if (r == NULL) {
        dev_err(&pdev->dev, "platform_get_resource failed\n");
        return -ENODEV;
    }

    if(request_mem_region(r->start, resource_size(r), "7iodev")!=NULL){
        printk(KERN_INFO "register 7iodev fail\n");
        return -EBUSY;
    }

    misc_register(&iodev_dev);

    data->reg_base = devm_ioremap(&pdev->dev, r->start, resource_size(r));
    if (data->reg_base == NULL) {
        dev_err(&pdev->dev, "unable to remap MMIO\n");
        return -ENOMEM;
    }

    error = devm_request_irq(&pdev->dev, platform_get_irq(pdev, 0), goldfish_iodev_handler, 0, "goldfish_iodev", NULL);
    init_waitqueue_head(&wait_q);
    spin_lock_init(&wait_q_lock);
    condition = 0;
    printk(KERN_ERR "iodev probe finished\n");

    return 0;
}
```

(Goldfish) goldfish_iodev_probe (2)

```
static const struct of_device_id goldfish_iodev_of_match[] = {
    { .compatible = "generic,goldfish_iodev", },
    {},
};
MODULE_DEVICE_TABLE(of, goldfish_iodev_of_match);

static const struct acpi_device_id goldfish_iodev_acpi_match[] = {
    { "GFSH0001", 0 },
    { },
};
MODULE_DEVICE_TABLE(acpi, goldfish_iodev_acpi_match);

static struct platform_driver goldfish_iodev_device = {
    .probe      = goldfish_iodev_probe,
    .remove     = goldfish_iodev_remove,
    .driver = {
        .name = "goldfish_iodev",
        .of_match_table = goldfish_iodev_of_match,
        .acpi_match_table = ACPI_PTR(goldfish_iodev_acpi_match),
    },
};
module_platform_driver(goldfish_iodev_device);
```

(Goldfish) Add New Interrupt

- We should call below function inside `device_probe()`
 - `static inline int devm_request_irq(struct device *dev, \`
`unsigned int irq, \`
`irq_handler_t handler, \`
`unsigned long irqflags, \`
`const char *devname, \`
`void *dev_id)`

*Example of calling devm_request_irq()

```
data->reg_base = devm_ioremap(&pdev->dev, r->start, resource_size(r));
if (data->reg_base == NULL) {
    dev_err(&pdev->dev, "unable to remap MMIO\n");
    return -ENOMEM;
}

error = devm_request_irq(&pdev->dev, platform_get_irq(pdev, 0), goldfish_io_dev_handler, 0, "goldfish iodev", NULL);
```

(Goldfish) Interrupt Handler

```
static irqreturn_t goldfish_iodev_handler(int irq, void *dev_id)
{
    spin_lock(&wait_q_lock);
    condition = 1;
    wake_up(&wait_q);
    spin_unlock(&wait_q_lock);
    return IRQ_HANDLED;
}
```

(Goldfish) Edit File Operations

- Exercise 1) You should fill the lines of 2, 3, 4 according to the comments

```
static ssize_t iodev_write (struct file *file, const char __user *buf, size_t size, loff_t *loff) {
    int i;
    uint32_t lpn;

    if (*loff & (PAGE_SIZE - 1) || size != PAGE_SIZE )
        return -EINVAL;

    lpn = (uint32_t)(*loff >> PAGE_SHIFT);
    copy_from_user((char*)kbuf, buf, size);

    /*(iodev_write)
    1. Read status register (until not DEV_READY)
    2. Write lpn to LBA_REG
    3. Write kbuf repeatedly (size of 4 bytes)f to BUF_REG
    4. Write WRITE_CMD to CMD_REG
    5. Wait for condition variable (condition)
    6. Increase offset
    7. Return size
    */
    while (readl(data->reg_base + STATUS_REG) != DEV_READY);    //1

    spin_lock_irq(&wait_q_lock);
    condition = 0;

    //2
    for (i = 0; i < PAGE_SIZE / sizeof(uint32_t); i++ ) {
        //3
    }

    //4

    wait_event_lock_irq(wait_q, condition, wait_q_lock);    //5
    spin_unlock_irq(&wait_q_lock);

    (*loff) += size;    //6

    return size;    //7
}
```

(QEMU) Add New Interrupt

- We should set irq count to 1 and irq number to 15

```
void goldfish_iodev_init(void)
{
    struct goldfish_iodev_state *s;

    s = (struct goldfish_iodev_state *)g_malloc0(sizeof(*s));
    s->dev.name = "goldfish_iodev";
    s->dev.base = 0;
    s->dev.size = 0x1000;
    s->dev.irq_count = 1;
    s->dev.irq = 15;
```

Only IRQ number 15 is left for new interrupt

(QEMU) goldfish_iodev_init

```
void goldfish_iodev_init(void)
{
    struct goldfish_iodev_state *s;

    s = (struct goldfish_iodev_state *)g_malloc0(sizeof(*s));
    s->dev.name = "goldfish_iodev";
    s->dev.base = 0;
    s->dev.size = 0x1000;
    s->dev.irq_count = 1;
    s->dev.irq = 15;
    s->iodev_fd = open("./iodev.disk", O_RDWR);

    goldfish_device_add(&s->dev, goldfish_iodev_readfn, goldfish_iodev_writefn, s);

    register_savevm(NULL,
                    "goldfish_iodev",
                    0,
                    IODEV_STATE_SAVE_VERSION,
                    goldfish_iodev_save,
                    goldfish_iodev_load,
                    s);
}
```

(QEMU) goldfish_iodev_write

```
static void goldfish_iodev_write(void* opaque, hwaddr offset, uint32_t value)
{
    struct goldfish_iodev_state* s = (struct goldfish_iodev_state*)opaque;
    int status;

    if ( offset < 0 ) {
        cpu_abort(cpu_single_env, "iodev_dev_read: Bad offset %" HWADDR_PRIx "\n", offset);
        return;
    }

    switch (offset) {
        case IODEV_CMD_REG:
            s->status = IODEV_BUSY;
            s->cmd = value;
            switch(s->cmd) {
                case READ_CMD:
                    status = goldfish_iodev_data_read();
                    break;
                case WRITE_CMD:
                    status = goldfish_iodev_data_write((void *)opaque);
                    break;
                default:
                    cpu_abort(cpu_single_env, "iodev_cmd: unsupported command %d\n", s->cmd);
                    return;
            }
            if ( status == 0 ) {
                s->status = IODEV_READY;
                goldfish_device_set_irq(&s->dev, 0, 1);
            } else {
                cpu_abort(cpu_single_env, "iodev_cmd: command:%d failed %d\n", s->cmd);
                return;
            }
            break;
        case IODEV_LBA_REG:
            s->lba = value;
            break;
        case IODEV_BUF_REG:
            s->iodev_buf[s->iodev_buf_pos] = value;
            s->iodev_buf_pos = (s->iodev_buf_pos+1) % 1024;
            break;
    };
}
```

(QEMU) goldfish_iodev_read

```
static uint32_t goldfish_iodev_read(void* opaque, hwaddr offset)
{
    struct goldfish_iodev_state* s = (struct goldfish_iodev_state*)opaque;

    if ( offset < 0 ) {
        cpu_abort(cpu_single_env, "iodev_dev_read: Bad offset %" HWADDR_PRIx "\n", offset);
        return 0;
    }

    switch (offset) {
        case IODEV_STATUS_REG:
            return s->status;
    };

    return 0;
}
```

Exercise

- Write Goldfish and QEMU code to “iodev” can emulate read and write requests from user
 - You can test your softdevice with `4k_write.c` from iCampus